

Institut für Architektur von Anwendungssystemen

Universität Stuttgart
Universitätsstraße 38
D - 70569 Stuttgart

Diplomarbeit Nr. ????

Bootstrapping Provisioning Engines for On-demand Provisioning in Cloud Environments

Lukas Reinfurt

Studiengang: Informatik

Prüfer: Jun.-Prof. Dr.-Ing. Dimka Karastoyanova

Betreuer: Dipl.-Inf., Dipl.-Wirt. Ing.(FH) Karolina Vukojevic

begonnen am: ????.2013

beendet am: ????.2014

CR-Klassifikation: ???

Abstract

Table of Contents

Abstract	2
Table of Contents	3
List of Figures	5
List of Tables	6
List of Abbreviations	8
1 Introduction	10
1.1 Task of this Diploma Thesis	10
1.2 Structure of this Document	10
2 Fundamentals	11
2.1 SimTech	11
2.2 Bootstrapping	12
2.3 Cloud Computing	13
2.4 Provisioning	14
3 Related Work	20
3.1 On-demand Provisioning for Simulation Workflows	20
3.2 Dynamic Provisioning of Web Services for Simulation Workflows	23
3.3 Architecture for Automatic Provisioning of Cloud Services	24
4 Requirements, Constraints and Design	27
4.1 Requirements	27
4.2 Constrains	28
4.3 Design	28
5 Bootstrapping Process	41
6 Implementation	42
6.1 Selecting Frameworks and Libraries	42

Table of Contents

7 Summary and Conclusion	46
Bibliography	47
Declaration of Authorship	50

List of Figures

2.1	TOSCA service template structure [based on 25].	17
2.2	OpenTOSCA architecture [based on 6].	19
3.1	Simplified overview of service binding strategies [based on 27].	21
3.2	Proposed architecture [based on 27].	22
3.3	Extended architecture with added provisioning manager [based on 21]	24
3.4	Simplified architecture for automatic provisioning of cloud services [based on 14]	25
4.1	Simplified overview of the single local component architecture	29
4.2	Simplified overview of the single remote component architecture	30
4.3	Simplified overview of the 2-tier architecture	32
4.4	Simplified overview of the cloned component architecture	33
4.5	Simplified overview of the 2-tier architecture with plugins	36
4.6	Simplified overview of the 2-tier architecture with a plugin repository	37
4.7	Simplified overview of the 2-tier architecture with asynchronous web service communication	38
4.8	Simplified overview of the 2-tier architecture with asynchronous web service and a messaging queue communication	39
4.9	Bootware internal communication with PubSub pattern.	40

List of Tables

6.1	Feature comparison of Java plugin frameworks	42
6.2	Feature comparison of Java PubSub libraries	44

List of Listings

List of Abbreviations

AMI	Amazon Machine Images, page 18
APT	Advanced Packaging Tool, page 26
AWS	Amazon Web Services, page 14
BPEL	Business Process Execution Language, page 17
BPMN	Business Process Model and Notation, page 17
CSAR	Cloud Service Archive, page 18
EC2	Elastic Compute Cloud, page 14
ESB	Enterprise Service Bus, page 23
JPF	Java Plugin Framework, page 43
JSPF	Java Simple Plugin Framework, page 43
NIST	National Institute of Standards and Technology, page 13
OASIS	Organization for the Advancement of Structured Information Standards, page 16
OSGi	Open Service Gateway initiative, page 43
PubSub	publish-subscribe pattern, page 39
RDC	Remote Desktop Connection, page 25
SDKs	Software Development Kits, page 14
SimTech	Simulation Technology, page 11
SLA	Service Level Agreement, page 13
SPI	Service Provider Interface, page 43
SSH	Secure Shell, page 25

List of Listings

SWfMS	Simulation Workflow Management System, page 12
TOSCA	Topology and Orchestration Specification for Cloud Applications, page 16
VMs	virtual machines, page 15
VNC	Virtual Network Computing, page 25

1 Introduction

Workflow technology and the service based computing paradigm were mostly used in a business context until now. But slowly they are extended to be used in other fields, such as eScience, where business centric assumptions that were previously true aren't reasonable anymore. One of these assumptions is that services should run continuously. This made sense in large enterprise where those services are used often. Science, on the other hand, often takes a more dynamic approach, where certain services, for example for simulation purposes, are only used at certain times. In those cases, it would make more sense to dynamically provision service only when they are needed.

1.1 Task of this Diploma Thesis

The task of this diploma thesis is to design a lightweight bootstrapping system that can kick off dynamic provisioning in cloud environments. It should be able to provision various provisioning engines in all kinds of cloud environments. The provisioning engines then handle the actual provisioning of required workflow systems and services. A managing component that keeps track of provisioned environments is also part of this system.

Support for different cloud environments and provisioning engines should be achieved through means of software engineering. A functioning prototype that supports Amazon as cloud environment and TOSCA as provisioning engine should be implemented.

1.2 Structure of this Document

...

2 Fundamentals

This chapter starts with a short overview of the SimTech project, of which this diploma thesis is a part of. Next, bootstrapping is defined in the context of this diploma thesis, since it can have various different meanings. Then, a short overview of the cloud landscape is presented, with focus on Amazons cloud offerings, since these are used in this diploma thesis. Finally, we take a look at provisioning solutions, in particular TOSCA, which is also used later in this thesis.

2.1 SimTech

Since 2005, the German federal and state government have been running the Excellence Initiative¹, which aims to promote cutting-edge research, thereby increasing the quality and international competitiveness of German universities. In three rounds of funding, universities have competed with project proposals in three areas: Institutional Strategies, Graduate Schools, and Clusters of Excellence. In total, the winners received approximately 3.3 billion euros since 2005, split up between 14 Institutional Strategies, 51 Graduate Schools, and 49 Clusters of Excellence [10, pp. 16-18].

Simulation Technology (SimTech) is one of the Clusters of Excellence that are funded by the Excellence Initiative. In a partnership between the University of Stuttgart, the German Aerospace Center, the Fraunhofer Institute for Manufacturing Engineering and Automation, and the Max Planck Institute for Intelligent Systems, it combines over 60 project from researchers in Engineering, Natural Science, and the Life and Social Sciences. The aim of SimTech is to improve existing simulation strategies and to create new simulation solutions [10, pp. 109].

In the SimTech project, seven individual research projects collaborate in nine different project networks, one of which is project network 8: *Integrated data management, workflow and visualization to enable an integrative systems science*. The goal of this project network is to

¹http://www.dfg.de/en/research_funding/programmes/excellence_initiative/index.html

build an easy-to-use infrastructure that supports scientists in their day to day work with simulations [19].

As part of this project, the SimTech Simulation Workflow Management System (SWfMS) was developed. It is a tool that enables scientists to easily create, manage and execute simulation workflows [11].

2.2 Bootstrapping

The term *to bootstrap sth.* appears to have originated in the early 19th century in the United States, where phrases like “pulling oneself up over a fence by the straps of one's boots” were used as a figure for an impossible task [28]. In the early 20th century the metaphor's sense shifted to suggest a possible task, where one improves one's situation by one's own efforts without help from others. An example of this can be found in James Joyce's *Ulysses* from 1922, where he writes about “others who had forced their way to the top from the lowest rung by the aid of their bootstraps” [13]. From there, the metaphor extended to the general meaning it has today, which is the act of starting a self-sustaining process that proceeds without help from the outside.

An early reference to bootstrapping in the context of computing dates back to 1953, describing the “bootstrapping technique” as follows: “Pushing the load button then causes one full word to be loaded into a memory address [...], after which the program control is directed to that memory address and the computer starts automatically. [...] [This] full word may, however, consist of two instructions of which one is a Copy instruction which can pull another full word [...], so that one can rapidly build up a program loop which is capable of loading the actual operating program.” [7, p. 1273].

The term bootstrapping is also used with a similar meaning in a business context, where it refers to the process of starting and sustaining a company without outside funding. The company is started with money from the founders which is used to develop a product that can be sold to customers. Once the business reaches profitability it is self-sufficient and can use the profits it generates to organically grow further [20].

In this diploma thesis, bootstrapping describes the process of starting a simple program, that without further help is able to start much more complex programs. These complex programs might require additional middleware, databases, or other components. During the bootstrapping process, all these dependencies will be set up automatically.

2.3 Cloud Computing

Cloud computing emerged in recent years as an alternative to traditional IT. Compared to traditional IT, it offers customers far more flexibility in terms of short term access to and scalability of resource, such as servers, databases, communication service, etc. This increased flexibility is the result of a combination of certain technologies and business models that, although having been around for a while individually, where combined only in recent years.

Since cloud computing is a relatively new phenomenon, there are many definitions of it scattered around. Vaquero et al. looked at over 20 of them and proposed the following definition:

“Clouds are a large pool of easily usable and accessible virtualized resources (such as hardware, development platforms and/or services). These resources can be dynamically reconfigured to adjust to a variable load (scale), allowing also for an optimum resource utilization. This pool of resources is typically exploited by a pay-per-use model in which guarantees are offered by the Infrastructure Provider by means of customized SLAs.”² [26]

The National Institute of Standards and Technology (NIST) also proposes a definition:

“Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction.” [17]

Today, there are many different cloud providers offering a huge selection of services. The range of providers spans from titans like Amazon³, Google⁴, Microsoft⁵, and IBM⁶ to small, focused providers like Heroku⁷ or Jelastic⁸ and even solutions to build own clouds, like OpenStack⁹. The next section shows Amazon's cloud services in more detail, since those will be used in this diploma thesis.

²Service Level Agreement (SLA): “An agreement that sets the expectations between the service provider and the customer and describes the products or services to be delivered, the single point of contact for end-user problems and the metrics by which the effectiveness of the process is monitored and approved.” [23]

³<http://aws.amazon.com>

⁴<https://cloud.google.com>

⁵<http://azure.microsoft.com>

⁶<http://www.ibm.com/cloud-computing>

⁷<https://www.heroku.com>

⁸<http://jelastic.com>

⁹<https://www.openstack.org>

2.3.1 Amazon Web Services

In 2006, Amazon started offering cloud resource under the umbrella of Amazon Web Services (AWS). Since then, their offerings steadily increased and do now comprise over 20 different products and services for computing, data storage, content delivery, analytics, deployment, management, and payment in the cloud [1].

The most relevant cloud offering for this thesis is Elastic Compute Cloud (EC2)¹⁰, Amazons IaaS offer. It allows customers to rent virtual server instances at an hourly rate. These servers are freely configurable, so virtually any software can be installed, making EC2 very versatile [2]. In addition to general purpose instances (M3), Amazon offers a wide selection of specialized instances, which are optimized for a specific purpose. These include instances optimized for computation performance (C3), memory-intensive applications (R3), or high storage instances (I2). For this thesis we will be using Amazons low cost micro instances (T1) [3].

Also of interest to this thesis is Elastic Beanstalk¹¹, Amazons PaaS offering. Customers can upload an application and Elastic Beanstalk takes care of deployment and scaling [4]. This makes it easier and quicker to use than EC2, but also less flexible. It could be used instead of a more manual approach with EC2.

Amazon offers multiple ways to interact with cloud resources. All AWS offerings can be controlled using the AWS Management Console, a web based management interface that allows customers to start, stop, and manage cloud resources on-demand. It also provides access to account and billing information [12]. Additionally, Amazon provides a command line interface, tools for Eclipse and Visual Studio, and Software Development Kits (SDKs) for several programming languages, including Java, .Net, Python, Ruby, and the Android and iOS platforms [24]. In this thesis, we will use the AWS Java SDK¹² to interact with Amazons cloud resources programmatically.

2.4 Provisioning

This section provides an overview of provisioning in general and describes some of the provisioning solution available today. It focuses in particular on TOSCA and Open TOSCA, since those are used in the prototypical implementation later on.

¹⁰<http://aws.amazon.com/ec2>

¹¹<http://aws.amazon.com/elasticbeanstalk>

¹²<https://aws.amazon.com/sdkforjava>

2.4.1 Overview

Setting up a complex distributed system with many different components scattered across multiple environments is a time-consuming task if done by hand. For this reason, many provisioning solutions have been created over the years to automate this process. They differ in some areas, which we will discuss later, but their core functionality is basically identical: They prepare all necessary resources for a certain task. Let's state that core functionality a little more precise with the following definition: Provisioning is, "in telecommunications, the setting in place and configuring of the hardware and software required to activate a telecommunications service for a customer; in many cases the hardware and software may already be in place and provisioning entails only configuration tasks" [8]. Since we are working in a cloud environment, we won't have to deal with hardware directly, but rather with virtual machines (VMs). So for us, provisioning means the creation and deletion of VMs in a cloud environment, as well as the installation, configuration, monitoring, running and stopping of software on these VMs [15].

There are many benefits to using an automated provisioning solution instead of doing things manually. The manual approach is limited by how much work a single person can do at any time, whereas an automatic approach is able to do much more work, in less time, and potentially in parallel. This makes it possible to manage huge infrastructures with very little resources, which can save money compared to a manual approach. Since every step that needs to be done to provision a system has to be written down, a detailed description of the whole provisioning process is created. This makes the whole process reproducible and less error prone, since the human factor is taken out. Parts of such a description can also be shared in a business or even between businesses, which makes the process of creating such a description potentially much more efficient.

The general process of working with provisioning software is very similar with all the different solutions. It can be described as a two step process. In step one, a description of the whole provisioning process has to be created using the tools provided by the particular solution. In general this involves creating a textual description in a certain format that is understood by the provisioning software that is to be used. In this description, we tell the software what virtual resources we need, what software should be installed on them and how everything should be configured. In step two, we pass this description to the provisioning software, which interprets and executes it.

Many different provisioning solutions exist today. Some cloud providers offer provisioning solutions that are particularly tailored to their cloud offerings, for example AWS CloudFor-

mation¹³, which can only be used to provision resources in the Amazon cloud. Then, there are more generally usable provisioning solutions that are not bound to any particular cloud provider. A few popular examples include Ansible¹⁴, Chef¹⁵, and Puppet¹⁶, and TOSCA, which we will discuss in detail later.

All these solutions differ in some form or another. A full feature comparison of different solutions is out of scope for this diploma thesis, but what follows is a short overview of some of the differences. As already mentioned, AWS CloudFront is bound to Amazons cloud platform, while the other solutions are not. Chef and Puppet both use a client server architecture, where each node that should be configured by them has to run a client program to communicate with a server node, whereas Ansible executes its command over SSH and therefor doesn't require additional software on the nodes that are configured. The solutions also differ in modularity and flexibility. While Ansible, Chef, Puppet and TOSCA are highly flexible and can be used in a fine grained modular fashion, this also makes them more complex to use, for example compared to AWS CloudFront.

2.4.2 TOSCA

Topology and Orchestration Specification for Cloud Applications (TOSCA) is a standard that is currently being worked on by the Organization for the Advancement of Structured Information Standards (OASIS)¹⁷. Its development is also supported by various industry partners, which include IBM, Cisco, SAP, HP and others. It's aim is to provide a language that can describe service components and their relations in a cloud environment independent fashion [25].

TOSCA defines an XML syntax, which describes services and their relations in a so called service templates. Figure 2.1 shows that a service template can be made of up of four distinct parts: Topology templates, orchestration plans, reusable entities, and artifact templates.

Topology templates, as seen on the left side of Figure 2.1, model the structure of a service as a directed graph. The vertices of the graph represent nodes, which are occurrences of a specific component, for example, an application server or a database. These nodes are defined by node types, or by other service templates. Node types are reusable entities, as shown in the top center of Figure 2.1. They define the properties of a component, as well as operations to manipulate a component, so called interfaces. Additionally, node types can

¹³<http://aws.amazon.com/cloudformation>

¹⁴<http://www.ansible.com>

¹⁵<http://www.getchef.com/chef>

¹⁶<http://puppetlabs.com/>

¹⁷<https://www.oasis-open.org/>

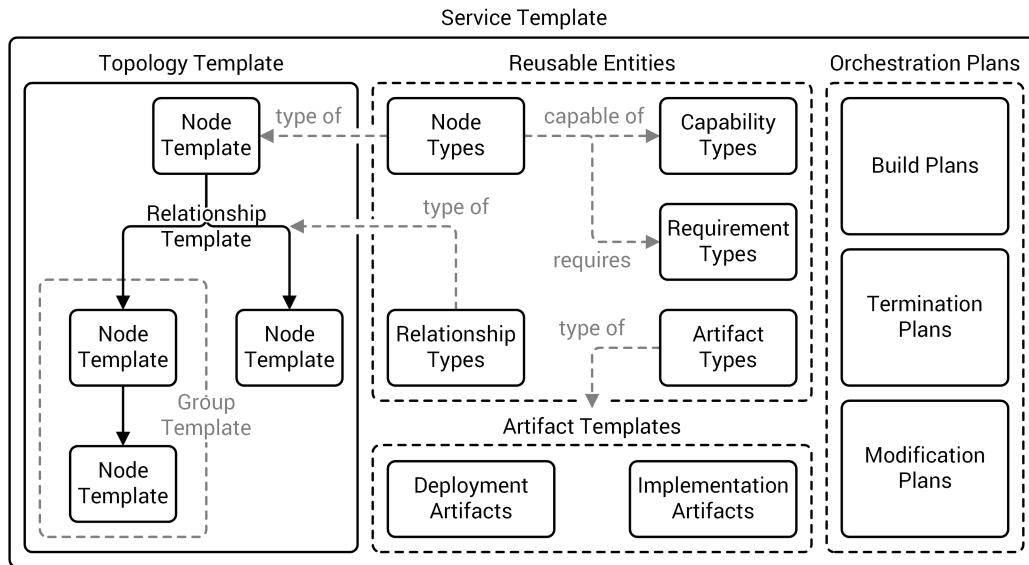


Figure 2.1: TOSCA service template structure [based on 25].

be annotated with requirements and capabilities. These, in turn, are defined by requirement and capability types, which also belong to the group of reusable entities. This allows for requirement and capability matching between different components. The edges of the graph represent connections between nodes, which are defined by relationship templates that specify the properties of the relation. An example for such a connection would be a node A, representing a web service, which is deployed on node B, an application server. Relationship types are also used to connect requirements and capabilities.

Orchestration plans, shown on the left of Figure 2.1, are used to manage the service that is defined by the service template. TOSCA distinguishes between three types of plans: Build plans, termination plans, and modification plans. Build plans describe, how instances of a service are created. Termination plans describe, how such a service is removed. Modification plans manage a service during its runtime. These plans consist of one or more tasks, i.e., an operation on a node (via an interface) or an external service call, and the order in which these tasks should be performed. They can be written in Business Process Model and Notation (BPMN) or Business Process Execution Language (BPEL), which are already existing languages to describe process models.

The bottom center of Figure 2.1 shows artifact templates, which represent artifact. Artifacts are things that can be executed directly (e.g.: scripts, archives) or indirectly (e.g.: URL, ports). TOSCA further distinguishes between two types of artifacts, namely deployment and implementation artifacts. Deployment artifacts materialize instances of a node and are used by a

build plan to create a service. An example for this is an Amazon Machine Images (AMI), which creates an Apache server once deployed in a VM. Implementation artifacts represent the interfaces of components. Here, an example would be a node that has an interface for starting the particular component described by the node. This interfaces could be implemented by an implementation artifact like a *.jar* file.

One or more TOSCA service templates are packaged, together with some meta data, into a Cloud Service Archive (CSAR), which is essentially a zip file that contains all files necessary to create and manage a service. CSAR files can then be executed in a TOSCA runtime environment, also called TOSCA container, to create the service described within.

2.4.3 OpenTOSCA

OpenTOSCA is a browser based open-source implementation of a TOSCA container, created at the IAAS at University Stuttgart, which supports the execution of TOSCA CSAR archives. Figure 2.2 shows the architecture of OpenTOSCA. Its functionality is realized in three main components, which are the Controller, the Implementation Artifact Engine, and the Plan Engine. After a CSAR is uploaded to OpenTOSCA it can be deployed in three steps. In the first step, the CSAR file is unpacked and its content is stored for further use. The TOSCA XML files are then loaded and processed by the Controller. The Controller in turn calls the Implementation Artifact Engine and the Plan Engine. The Implementation Artifact Engine knows how to deploy and store the provided implementation artifacts via plugins. Plans are then run by the Plan Engine, which also uses plugins to support different plan formats. OpenTOSCA also offers two APIs, the Container API and the Plan Portability API. The Container API can be used to access the functionality provided by the container from outside and to provide additional interfaces to the container, like the already existing admin UI, self-service portal, or modeling tool. The Plan Portability API is used by plans to access topology and instance information [6].




Figure 2.2: OpenTOSCA architecture [based on 6].

3 Related Work

This chapter summarizes related work of other authors that is of interest to this thesis. First, we present the paper that laid the foundation for this diploma thesis. Then we take a look at another diploma thesis which expanded the ideas presented in the first paper. We also take a look at a paper that presents work similar to this diploma thesis.

3.1 On-demand Provisioning for Simulation Workflows

Vukojevic-Haupt, Karastoyanova, and Leymann identify requirements that need to be addressed to make the current approach for scientific workflows used by the SimTech SWfMS more suitable for scientific simulation work. These requirements are: Dynamic allocation as well as release of computing resources, on-demand provisioning and deprovisioning of workflow middleware and infrastructure, and dynamic deployment and undeployment of simulation services and their software stacks. To fulfill these requirements, they propose a new service binding strategy that supports dynamic service deployment, an approach for dynamic provisioning and deprovisioning of workflow middleware, an architecture that is capable of these dynamic deployment and provisioning operations, and, as part of this architecture, the bootware - the subject of this diploma thesis - that kicks off these dynamic processes [27].

The new service binding strategy is necessary, because existing static and dynamic binding strategies, as shown on the left and in the center of Figure 3.1, rely on services that are always running, or, as in the case of dynamic binding with service deployment, only dynamically deploy the service, but not its middleware and infrastructure. The new service binding strategy, shown on the right of Figure 3.1, called *dynamic binding with software stack provisioning*, is similar to the already existing dynamic binding with service deployment strategy, but adds the dynamic provisioning of the middleware and infrastructure required by the service [27].

Their approach for dynamic provisioning and deprovisioning of workflow middleware is separated into six steps. The first step is to model and start the execution of a simulation

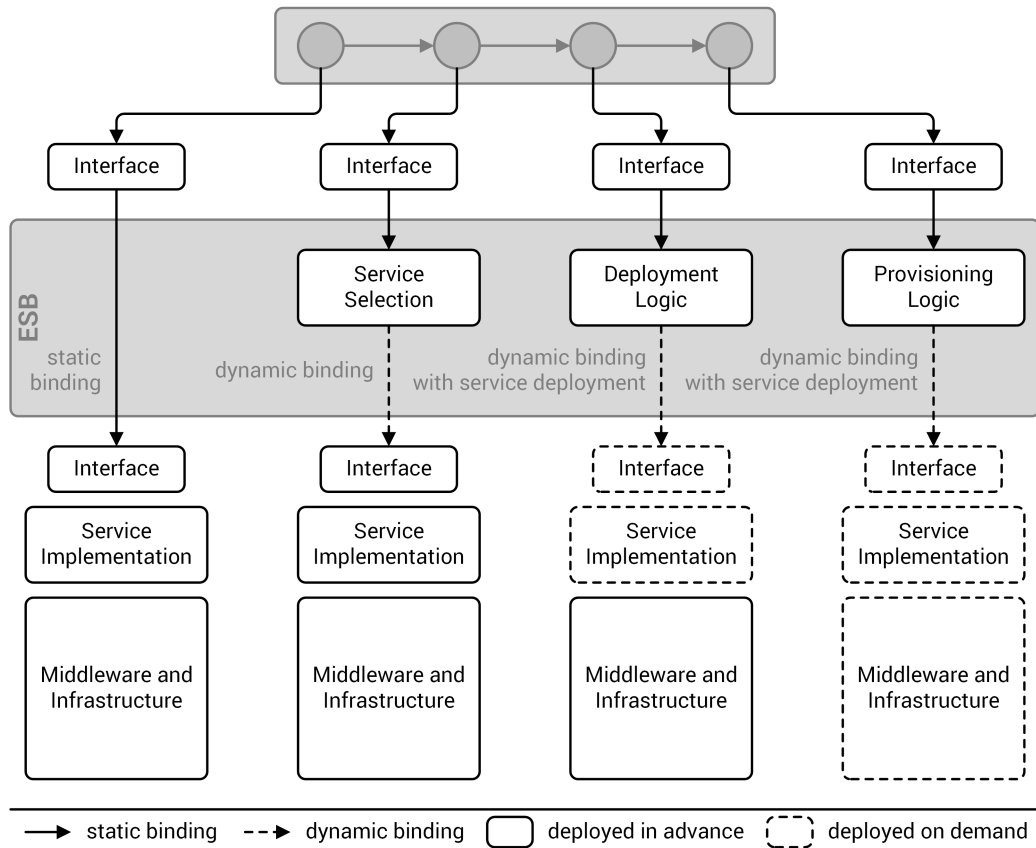


Figure 3.1: Simplified overview of service binding strategies [based on 27].

workflow. For this, the Modeler component shown on the left of Figure 3.2 is used. In the second step, the middleware for executing the workflow, e.g. the SWfMS, and its underlying infrastructure are provisioned. This is shown in Figure 3.2 as deployment step one and two. In deployment step one, the bootware deploys a provisioning engine into a cloud environment. In deployment step two, this provisioning engine is used to deploy the actual middleware. Now, the workflow can be deployed on this middleware, which is step three. In step four, an instance of this workflow is executed. During this execution, some external service might be required that are not yet available. The ESB determines this by checking the service registry. If the requested service isn't available, the ESB tells the provisioning engine to provision this service. The on-demand provisioning of services is step five, which corresponds to deployment step three in Figure 3.2. The service is also deprovisioned if it is no longer needed. The final step is to deprovision the workflow model and the workflow execution middleware after the execution of the workflow instance is finished [27].




Figure 3.2: Proposed architecture [based on 27].

The architecture they present can be separated into a local part and a cloud part, as well as different phases. Figure 3.2 shows that the only local components are the modeler and the bootware, while all other components are hosted in the cloud. In the modeling phase, a scientist uses local modeling and monitoring tools in combination with cloud hosted repositories and registries to create a workflow. These components are always running. When he deploys the workflow, the local bootware component kicks off the on demand provisioning process and therefore the second phase, called middleware runtime phase. In this phase, the bootloader deploys a provisioning engine in the cloud, which in turn deploys the workflow middleware. Once the middleware is up and running, the workflow can be executed. During the execution, the ESB receives service calls from the workflow engine. Services, that are not running at this time can then be provisioned by the provisioning engine. This takes place in the third phase, the service runtime phase. The bars at the bottom of Figure 3.2 show, which components are active during which phase [27].

3.2 Dynamic Provisioning of Web Services for Simulation Workflows

Schneider found some problems with the architecture proposed in section 3.1. The original architecture assumes that only one provisioning engine is used at a time. It neglects situations where services might require another (or multiple other) provisioning engines, because their provisioning descriptions aren't available in a format that the currently used provisioning engine understands. It also assumes that the Enterprise Service Bus (ESB) communicates directly with this provisioning engine to deploy and undeploy other services. This implicates that the ESB understands all manner of interfaces provided by various provisioning engines.

It further assumes that every provisioning engine knows how to communicate with the service repository to get the information and resources it needs to provision a service. While this might be true for some provisioning engines, it's certainly not true for all of them. This problem is further amplified because there are no standards defined for such a service repository.

Another assumption of the original architecture is that a provisioning engine always understands the format of the service packages provided by the service repository. Different provisioning engines use different formats which are in general not compatible. If provisioning engines would all use a standardized format (like CSAR), this would not be a problem, but that isn't the case.

Schneider further refines the previously shown middleware architecture by adding a provisioning manager as intermediary between the ESB and the provisioning engines [21]. This addition improves the original architecture in three aspects.

The ESB can now use the stable interface of the provisioning manager to trigger provisioning engines instead of calling those provisioning engines directly. The provisioning manager handles the differences between the provisioning engines. This makes it also possible to use multiple different provisioning engines during one workflow execution. The provisioning manager also handles the communication with the service repository or possibly multiple service repositories for different provisioning engines. He can provide information to a particular provisioning engine if it cannot get the information it needs from the service repository on its own. The provisioning manager could also translate different service distribution formats so that provisioning engines could be used with formats that they don't support

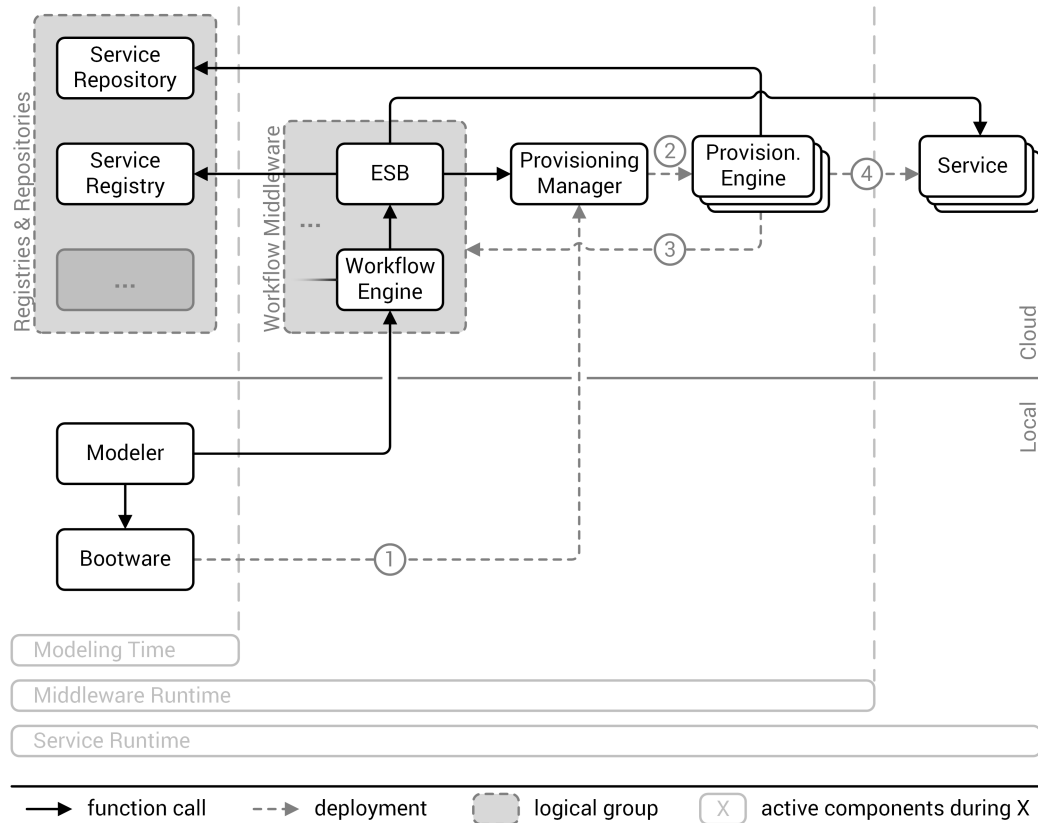


Figure 3.3: Extended architecture with added provisioning manager [based on 21]

3.3 Architecture for Automatic Provisioning of Cloud Services

Kirschnick et al. present an extensible architecture for automatic provisioning of cloud infrastructure and services at different cloud providers. They define a cloud service as a number of software components, their configuration, and the cloud infrastructure they are running on. Therefore, to provision a cloud service, one must first provision the cloud infrastructure, then install the software components and finally configure them. For this process they designed a so called service orchestrator [14].

Figure 3.4 shows a simplified overview of the service orchestrator architecture. Users can submit service models via a service API, shown at the bottom of Figure 3.4. The service model describes the topology of a cloud service, the components it consists of and their configuration. This model is used by the service orchestrator to provision new cloud services




Figure 3.4: Simplified architecture for automatic provisioning of cloud services [based on 14]

and to trigger reconfiguration and topology changes of existing services [14].

The service orchestrator is build in two layers. The orchestration layer, shown at the bottom of Figure 3.4, picks up the service model and delegates the different steps that are necessary to provision the cloud service to the abstraction layer. The deployment steps will be executed in parallel wherever dependencies allow it. The abstraction layer, shown in the middle of Figure 3.4, provides abstract methods to handle the management, installation, configuration, and starting of software. It consists of five different managers for infrastructure, packages, applications, configuration, and VM connections [14].

The infrastructure manager abstracts away different APIs of cloud providers for provisioning cloud infrastructure. It is used to create new VMs at a specific cloud provider, shown at the top of Figure 3.4. These can then be used by the other managers to create the requested cloud service. The VM connection manager offers a unified interface for different communication methods, for example Secure Shell (SSH), Remote Desktop Connection (RDC), Virtual Network Computing (VNC), and Telnet. The package, configuration, and application managers use the interface provided by the connection manager to send their commands to

3 *Related Work*

the VMs at the cloud provider [14].

The package manager can install software packages in different environments, for example with existing package managers like Advanced Packaging Tool (APT), or directly on the file system. The configuration manager offers a unified interface for component specific configuration methods, such as file templates or configuration APIs. The application manager manages the state of software components, i.e. starting and stopping them. Together, these managers can install, configure, and manage the requested cloud service on the VMs provided by the cloud provider [14].

4 Requirements, Constraints and Design

In this chapter we will talk about the design of the bootware component and the requirements and constraints that shape it. We begin with the requirements, which were explicitly given at the beginning of the thesis. Next, we describe additional constraints which we added to limit the scope of the work. Then, we explain the design that formed as a result of the limitations that we had.

4.1 Requirements

The main goal of this thesis is to lay a ground stone by creating the core design of the bootware component. It was clear from the beginning that, because of the limited time available, not every feature that might be necessary for the full operation can be fully implemented. Instead, the foundation we develop here should keep future needs in mind and make it simple to extend when needed. It is therefore a core requirement to keep the bootware component relatively generic and make it extensible where necessary.

It should be extensible in two key areas, namely the support for different cloud providers and for different provisioning engines. For this thesis, Amazon is the only cloud provider that has to be supported, but it has to be possible to add others in the future. Concerning provisioning engines, only OpenTOSCA has to be supported for now, but again with the possibility to add more in the future.

It is also important that the bootware is easy to use. In fact, it should be practically invisible whenever possible. It should hook into the already existing process of deploying a workflow without adding unnecessary interaction steps when possible. However, It can't be hidden completely, because the user has to specify a cloud provider and the corresponding log-in credentials some where. The user should also get some feedback about the progress of the deployment, because this process might take some time and might seem unresponsive without sufficient feedback.

A further requirement is that the bootware component should be relatively lightweight and open standards should be used where possible.

4.2 Constrains

The Bootware could theoretically be written in any major programming language but we limit our selves to Java. The reason for this is that all the other SimTech components are written in Java, so by also using Java we fit nicely into this already existing ecosystem. Additionally, for things like Eclipse integration we would have to use Java anyway. We also have to keep in mind that the bootware component will not be finished with this thesis. Other people will have to extend it in the future and since Java is common in general, as well as in the SimTech project, it makes sense to use it instead of another, maybe more obscure programming language.

We can further narrow our use of Java by limiting us to Java 1.6. This also has to do with the already existing parts of the SimTech project, that are geared towards this version as well. Using another version of Java could lead to unforeseen incompatibilities.

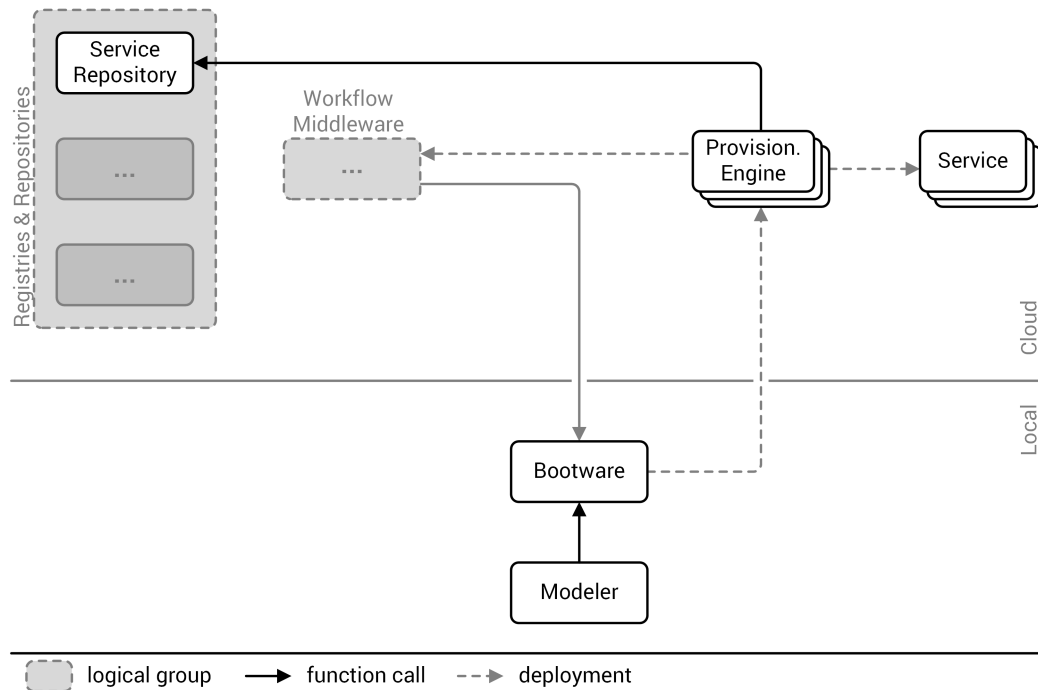
In the next chapter we will also introduce additional constraints that became necessary during the design process and will therefore be explained at the appropriate times.

4.3 Design

In this section we will develop the design of the bootware. This design is held intentionally abstract. Specific implementation details will be selected and described in section 6.1.

4.3.1 Component Division

The first question that has to be answered is whether or not it makes sense to divide the bootware component into two or more cooperating components, i.e. a server client division or a similar structure. *As described in section 3.1 on page 20, the proposed architecture initially only envisioned one bootware component. This architecture was expanded with the introduction of the provisioning manager, as described in section 3.2 on page 23.* We will now discuss the viability of these architectures.

Single Local Component**Figure 4.1:** Simplified overview of the single local component architecture

First we consider the simplest case: A single local component. In this scenario, all provisioning processes are initiated from a component installed locally on the users machine, alongside or as part of the workflow modeler.

The advantages of this architecture lie in its simplicity. Only one component has to be created and managed. We wouldn't have to deal with any cloud environments and each user would have his own personal instance, so multi-tenancy wouldn't be an issue. There is no possible overlap in functionality, as it would be the case in a 2-tier architecture (see section 4.3.1) and communication between multiple bootware components doesn't have to be considered.

The disadvantages are caused by the component being local. Since all the functionality is concentrated in one component, it can become quite large and complicated, which is one thing that should be avoided according to the requirements. A much bigger problem however is the remote communication happening in this scenario. All calls to the bootware component from the ESB of the workflow middleware would leave the remote environment. Also, all calls from the bootware component to the provisioning engines would enter the remote environment. This type of split communication can be costly and slow, as shown by Li et al.

They compared public cloud providers and measured that intra-datacenter communication can be two to three times faster and also cheaper (often free) compared to inter-datacenter communication [16].

Single Remote Component

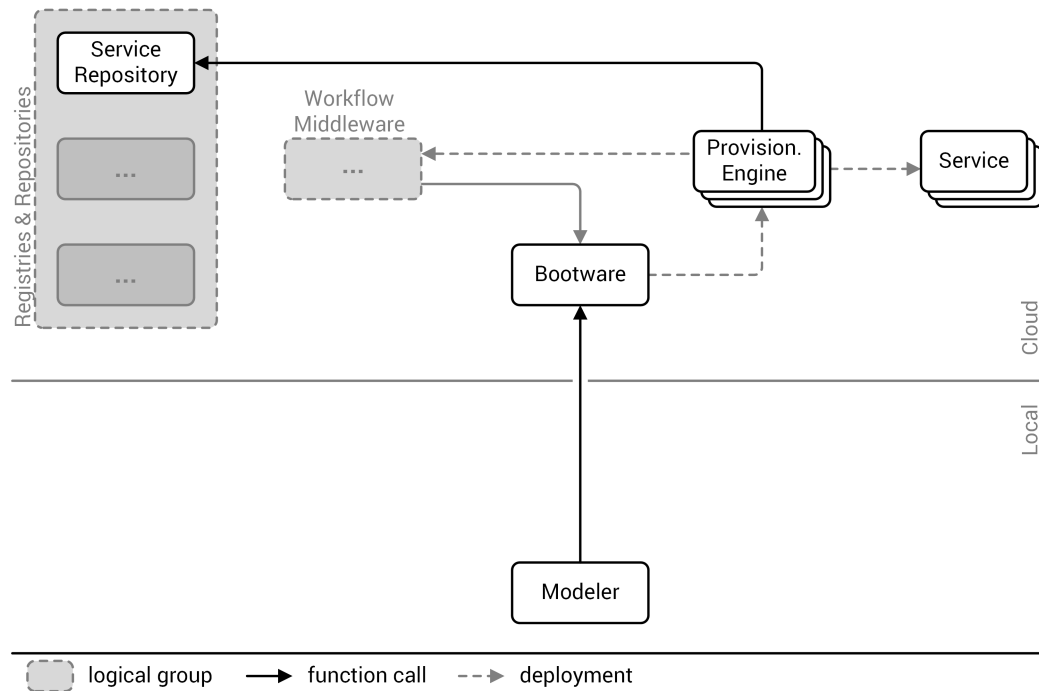


Figure 4.2: Simplified overview of the single remote component architecture

The next obvious choice is to put the single bootware component into a remote environment, where the disadvantages of local to remote communication would disappear. However, this creates new problems.

Since there aren't any additional components in this scenario that could manage the life-cycle of the remote bootware, the user would have to manage it by hand, which leads to two possibilities. Either the user provisions the bootware once in some cloud environment and then keep this one instance running or she provision it once she needs it and deprovisions it when she is done.

In the first case the user would only have to provision the bootware once, but this creates a new problem: The user doesn't know where exactly to put the bootware component. Since

one requirement is that multiple cloud environments should be supported, it is possible that the bootware component is not located anywhere near the cloud environment where it should provision further components. The communication problem of the single local bootware component can still occur in these cases.

Another problem is that the bootware would be running all the time, even if the user doesn't need it, which would increase costs. This problem could be reduced if this bootware instance is shared with others to assure a more balanced load. But then the user would have to manage some sort of load balancing and the bootware component would have to support multi-tenancy or be stateless to be able to cope with potential high usage spikes. This would further complicate the design and implementation of the bootware component and possibly increase the running costs.

In the second case the user would provision the bootware whenever she needs it. Now the user would be able to pick a cloud environment that is close to the other components that she plans to provision later. This eliminates the two major problems of the first case but increases the effort that the user has to put into a task that she shouldn't have to do in the first place. Life-cycle management of the bootware should be automated completely and hidden away from the user. Therefore, this scenario is not appropriate for our case.

2-Tier Architecture

Next, we take a look at a 2-tier architecture, where the bootware is divided into two components. On the local side we have a small and simple component which has mainly one function: To provision the larger second part of the bootware in a remote environment near to the environment, where other components will be provisioned later.

This eliminates the problems of a single local or remote bootware component. The user no longer has to be involved in the management of the remote bootware since the local bootware handles all that. Since we provision the remote bootware on demand we now also can position the remote component close to other remote components to minimize local/remote communication and the problems resulting of it. We can now keep the local part as simple as possible and make the remote part as complicated as it has to be and since we provision the remote bootware only for one user we don't have to worry about multi-tenancy.

But we also introduce new problems. For one, we now have duplicate functionality between the two components. Both components have to know how to provision a component into multiple cloud environments. The local component has to be able to put its remote counterpart into any cloud environment. The remote component has to be able to provision other components




Figure 4.3: Simplified overview of the 2-tier architecture

into the same environment in which it runs (ideally, to minimize costs). Since itself can be located in any cloud environment, it has to be able to do this in any cloud environment. Independent from this, it has to be able to provision to any environment that the user/service package chooses. But this problem can be solved by using a plugin architecture, which allows both components to use the same plugins. We discuss plugins in detail in subsection 4.3.2.

A second problem which we can't avoid but can solve is the communication which is now necessary between the different parts of the bootware. More on this in subsection 4.3.3

Cloning

This architecture can be seen as an alternative form of the 2-tier architecture described in section 4.3.1. In this case there are also two parts working together and the remote part does most of the work. However, the local and the remote components are identical. Instead of provisioning a bigger component in a remote environment, the local component clones itself. Compared to the 2-tier architecture described before, this has the advantage that only one component has to be designed and implemented and that function duplication is not




Figure 4.4: Simplified overview of the cloned component architecture

an issue. The disadvantage would be that the local bootware would be exactly as complex as the remote bootware and would contain functionality that it wouldn't require for local operation (e.g. a web service interface). However, since we want to keep the whole bootware, including the remote part, fairly lightweight, it's highly unlikely that the complexity of the remote bootware will reach such heights that it could not be run on an average local machine. In this case, the advantage of only having to design and implement one component seems to outweigh the disadvantage of a slightly more complex local component (compared to the 2-tier variant). Of course, this architecture makes only sense if the functionality of the two separate components in the 2-tier architecture turns out to be mostly identical. Therefore we can't decide yet if this architecture should be used.

Decision

Of the four alternative presented here, alternative three - the 2-tier architecture - makes the most sense. Therefore it is selected as the alternative of choice and used for further discussion. We do however retain the option to transform it into alternative four if we discover

that both components share much of same functionality, but this can only be judged at a later stage, when we know exactly how the internal functionality of the bootware will work.

4.3.2 Extensibility

The requirements for the bootware component state that support for different cloud environments and provisioning engines should be achieved through means of software engineering. This requirement is intentionally vague to allow to select a fitting extension mechanism during the design process. In this section we will take a look at different extension mechanisms for Java and pick the one that suits our need best.

extension definieren?

Extension Mechanisms

The simplest way to fulfill the extensibility requirement would be to create a set of interface and abstract classes to define the interfaces and basic functionality that are necessary to work with different cloud environments and provisioning engines. These interfaces and abstract classes would then be implemented separately to support different scenarios and would be compiled, together with the rest of the application, into one executable. At runtime, a suitable implementation would be selected and used to execute the specific functionality required at this time.

This extension mechanism is simple, but restricted by its static nature. The entire executable has to be recompiled if any implementations are changed or added. This may not be a problem if the set of possible extensions that have to be supported is limited and known at the time of implementation or if it changes rarely. If the set of necessary extensions is unknown or changing from time to time, implementing new, or changing existing extensions, can get cumbersome, since a new version of the whole software has to be released each time. Additionally, if a user wanted to add an extension for a specific use case, he would have to get the source code of the bootware, add his extension and then compile the whole bootware. It would be far better if extensions could be implemented separately from the core bootware component and added and removed at will.

A more flexible architecture is needed, for example a plugin architecture. Interfaces for the extension points still exist but the extension are no longer part of the main bootware compo-

nent. They are compiled separately into plugins that can be loaded into the main bootware component on the fly. There are several possibilities to realize such an architecture.

It is certainly possible to implement a plugin framework from scratch. An advantage of this approach would be that the design of the plugin architecture could be tailored to our use case and would be as simple or complex as needed. But there are also several disadvantages. For one, we would reinvent the wheel, since multiple such frameworks already exist. It would also shift resources away from the actual goal of this thesis, which is designing the bootware component. Furthermore it would require a deep understanding of the language used for the implementation (in this case Java), which is not necessarily given. Therefore it seems more reasonable to use one of the already existing plugin frameworks. Which one exactly will be determined later in subsection 6.1.1.

Plugin Repository

Now that we have introduced plugins we face new problems. Figure 4.5 shows the current architecture, where both bootware components use their own plugins. If a plugin is added or updated, the user has to manually copy this plugin to the right folder of one or both of the bootware components. Furthermore, if both components use the same plugins, which they will (for example plugins for different cloud providers), we will have duplicate plugins scattered around. This is inefficient, probably annoying for the user and can possibly cause errors if plugins get out of sync.

To remedy this situation we introduce a central plugin repository, as shown in Figure 4.6. This repository holds all plugins of both components so it eliminates duplicate plugins. If plugins are added or modified it has only to be done in one place. Plugin synchronization can happen automatically when the bootware components start, so that the user is no longer involved in plugin management. The repository also enables easy plugin sharing, which was cumbersome earlier.

While a central plugin repository is a sensible addition to the proposed bootware architecture, its design and implementation are out of scope of this thesis.

4.3.3 External Communication

Since we will use a 2-tiered approach for the bootware component, we now have to decide, how the communication between the two components will work. There are several factors




Figure 4.5: Simplified overview of the 2-tier architecture with plugins

that impact this decision.

Communication between the components should be as simple as possible, but has to support some critical features.

Since the provisioning processes kicked off by the bootware can potentially take a lot of time to finish (in the range of minutes to hours), asynchronous communication should be used between the components to avoid timeouts and blocking resources.

For the same reason, there should be some mechanism to get feedback on the current status during a long running provisioning process.

There will be other components that have to call the bootware to start some provisioning process, so communication has to be open for other parties.

The communication with the bootware components will contain sensitive data, for example login information for cloud providers that the bootware needs to do its job, but has to be provided from the outside on a call to call basis. This information has to be transported securely to prevent malicious or fraudulent attacks. The selected communication method

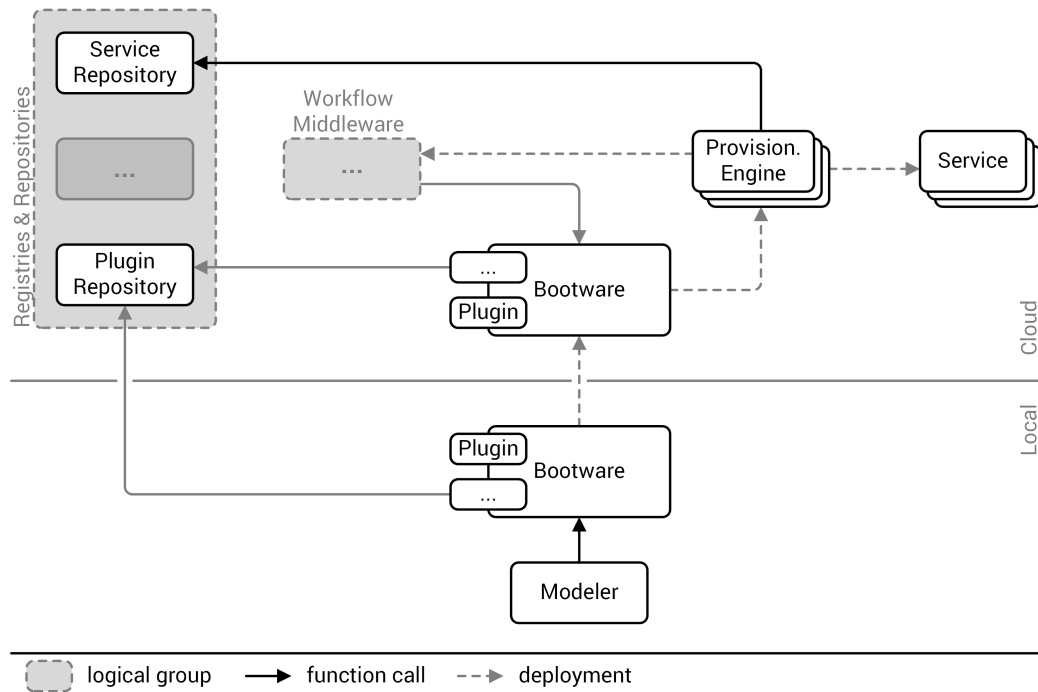


Figure 4.6: Simplified overview of the 2-tier architecture with a plugin repository

therefore has to support some sort of security mechanisms, ideally end-to-end encryption. While these security mechanisms will not be used in this thesis due to time constraints, selecting the right communication method is still critical for future development.

REST, RMI (restricted to java clients), EJB, RPC

Since this whole project is concerned with services, it *liegt nahe* to use web service calls and returns as main communication method, internally as well as externally. This way we offer a consistent interface for all components to use. Figure 4.7 shows the addition of asynchronous web service call and return communication to the proposed architecture. Using a web service with soap messaging also gives us access to Web Service Security Framework, which supports end-to-end encryption.

Web Services also support asynchronous communication so that long running provisioning processes won't pose a problem. We do however still need information during those long running processes to give the user some feedback. This can't be accomplished by this simple request/response pattern. For this, a secondary communication mechanism which supports sending multiple feedback messages has to be used.

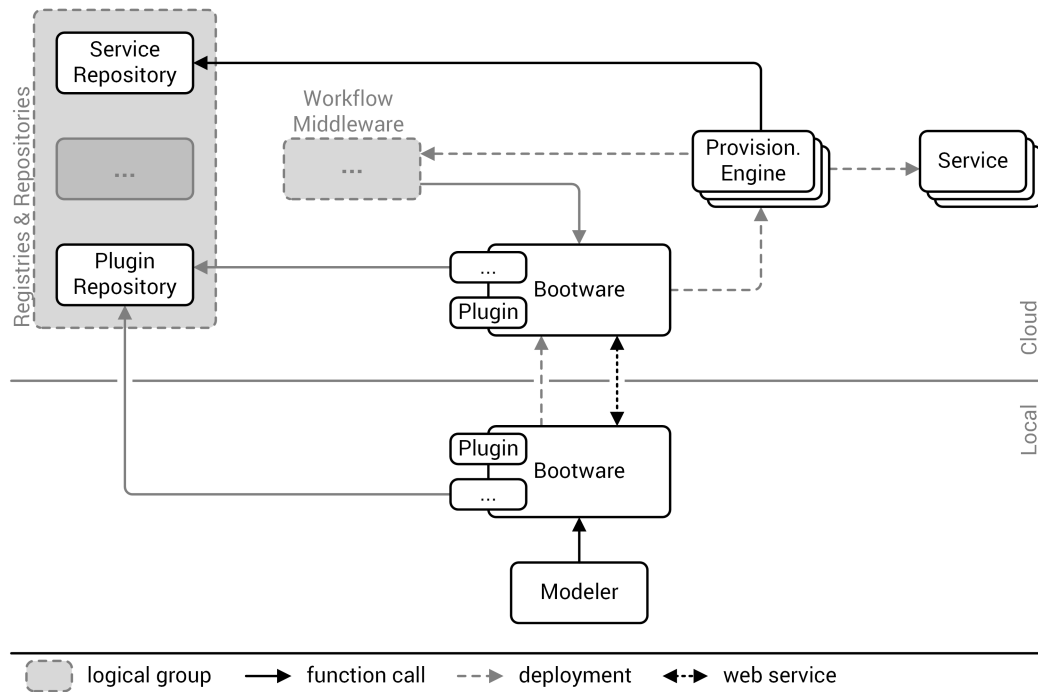


Figure 4.7: Simplified overview of the 2-tier architecture with asynchronous web service communication

Since it is not necessary for the successful use of the bootware it would make sense to implement this secondary communication mechanism as a plugin. This would allow us to add arbitrary communication plugins to the bootware depending on future needs.

One natural choice for this is the PubSub pattern implemented by some messaging middleware. Using this pattern, the remote bootware component pushes messages to a message queue to which the local bootware component (and other components if needs be) can subscribe to receive future messages. Figure 4.8 shows an additional message queue.

4.3.4 Internal Communication

We also have to consider internal communication between the bootware core and plugins, and possibly also in between plugins. Ideally, every plugin will be able to react to events from the bootware. These event could be triggered by the bootware core or by any plugin, but plugins should be completely independent from each other. Since a plugin doesn't know about other plugins, it can't listen for events at other plugins directly. The only known constant to a plugin

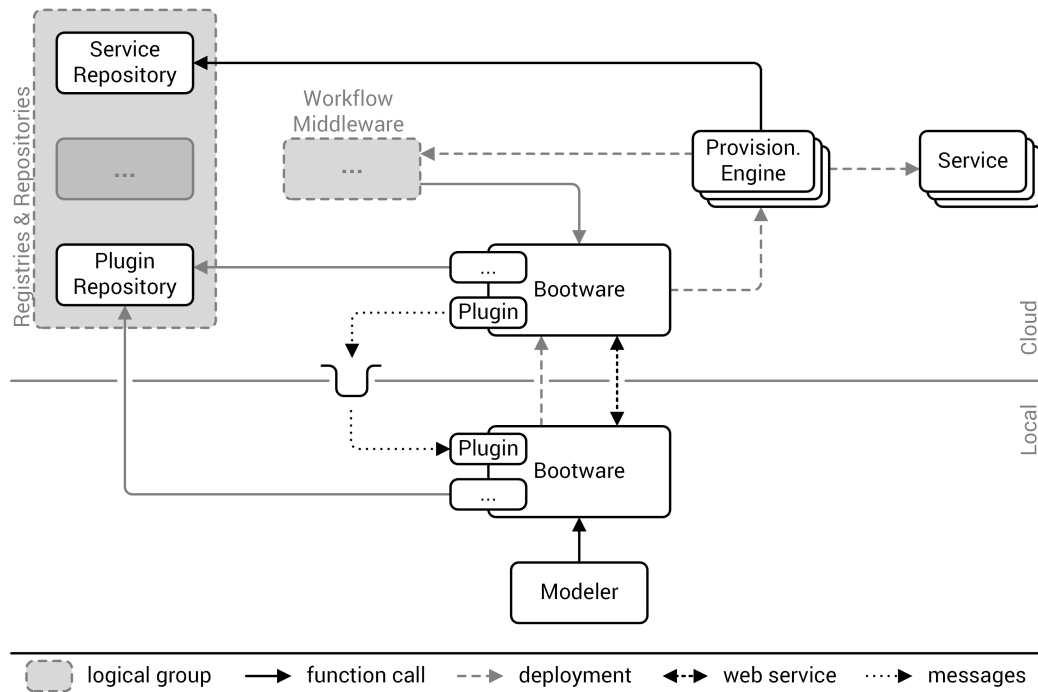


Figure 4.8: Simplified overview of the 2-tier architecture with asynchronous web service and a messaging queue communication

is the bootware core. Therefore we need a communication mechanism which allows for loosely coupled communication between the bootware core and the plugins, where plugins can register their interest for certain events with the core and also publish their own events to the core for other plugins to consume. This essentially describes the publish-subscribe pattern [9].

PubSub

The publish-subscribe pattern (PubSub) is a messaging pattern that consists of three types of participant: A event bus (or message broker), publishers, and subscribers. The event bus sits at the center of the communication. He receives messages from publishers and distributes them to all subscribers that have voiced their interest in messages of a certain type by subscribing at the event bus [9].

Using this pattern in our bootware component, we would create an event bus at the bootware

core and plugins, as well as other parts of the core, could subscribe at this event bus and also publish messages through this event bus.

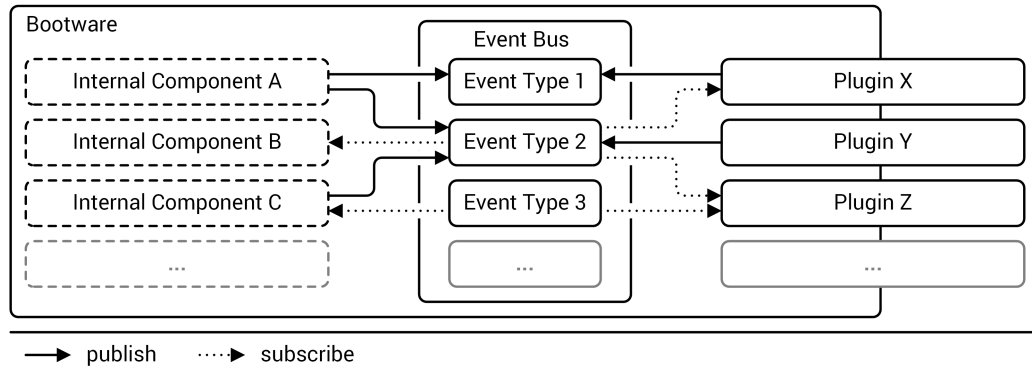


Figure 4.9: Bootware internal communication with PubSub pattern.

5 Bootstrapping Process

This chapter describes the bootstrapping process in its entirety.

6 Implementation

6.1 Selecting Frameworks and Libraries

6.1.1 Plugin Frameworks

		Plugin Frameworks			
		SPI ¹	JSPF ²	JPF ³	OSGi ⁴
<i>functional</i>	Dynamic Loading	x	x	✓	✓
	Security	x	x	x	✓
<i>non-functional</i>	Used in SimTech	x	x	x	✓
	Standard	x	x	x	✓
	Complexity	x	low	medium	high
	Popularity	x	low	low	high
	Active Development	x	x	x	✓

Table 6.1: Feature comparison of Java plugin frameworks

All of the frameworks that we compare here offer the basic functionality that we need to extend the core bootloader component, i.e. the developer defines interfaces that then are implemented by one or more plugins. These plugins are compiled separately from the main component and are then packaged in *.jar* files for distribution. These packages are loaded during runtime and provide the implementation for the specific interface they implement. There are however some advanced functional differences and some non-functional differences that will be considered here.

Dynamic loading allows us to load and replace plugins during runtime, without completely

¹<http://docs.oracle.com/javase/6/docs/api/java/util/ServiceLoader.html>

²<https://code.google.com/p/jspf>

³<http://jpf.sourceforge.net>

⁴<http://www.osgi.org>

6 Implementation

restarting the application. While we don't know for certain if dynamic loading is needed in our case, it's one of the advanced features that might be nice to have in the future.

Security is a must have feature but is out of the scope of this thesis. Consider the following scenario: The bootware component is used by multiple separate users who can share plugins using a plugin repository. Without security features, a malicious user could upload a plugin to this repository which, in theory, could contain any code. Therefore it's important to select the right framework now, so that security features can be implemented in the future.

Some non-functional features should also be considered, such as complexity, popularity, and if the framework is still in active development.

Java provides an integrated extension mechanism called Service Provider Interface (SPI), which is a little more advanced than the manual extension mechanism described above. It is also based on a set of interfaces and abstract classes that have to be implemented by an extension. In the case of SPI, these interfaces and abstract classes are called services and a specific implementation of such a service is called service provider. However, unlike in the above approach, specific implementations are loaded from .jar files in specific directories or in the class path. These .jar files also include metadata to identify the different service providers [22].

Java Simple Plugin Framework (JSPF)² is a plugin framework build for small to medium sized projects. Its main focus is simplicity. Therefore it does not support many of the advanced features like dynamic loading or security that other solution support. The author explicitly states that it is not intended to replace JPF or OSGi [5].

Java Plugin Framework (JPF)³ is an open-source plugin framework. Compared to JSPF it supports some advanced features like dynamic loading of plugins during runtime. It is also more popular than JSPF. However, the last version was released in 2007. This is not necessarily bad but might show that there will be no future development of this framework.

Open Service Gateway initiative (OSGi)⁴ is a plugin framework standard developed by OSGi Alliance. It provides a general-purpose Java framework that supports the deployment of extensible bundles [18].

Decision

6.1.2 Web Service

6.1.3 PubSub Libraries

Many of the well know Messaging Middlewares offer support for PubSub, for example ActiveMQ⁵, RabbitMQ⁶, and ZeroMQ⁷. But, since we are looking for an internal communication mechanism only, all of these solutions are somewhat overpowered. We don't have to worry about network problems, so we don't need guaranteed delivery or message queuing capabilities. We also don't need persistence or transactional capabilities. We don't have to handle millions of subscribers or events, so high scalability isn't a concern. We don't even necessarily need asynchronous communication. Instead, we need a lightweight in-memory solution. Therefore we will ignore the middleware heavyweights and look for smaller, more light weight PubSub libraries.

		PubSub Libraries				
		MBassador ⁸	Guava Event Bus ⁹	Simple Java Event Bus ¹⁰	EventBus ¹¹	Mycila PubSub ¹²
Features	Synchronous Invocation	✓	✓	✓	✓	✓
	Asynchronous Invocation	✓	✓	✓	✓	✓
	Filtering	✓	✗	✓	✓	✗
	Event Hierarchy	✓	✓	?	?	✓
	Prioritization	✓	✗	✗	✓	✗
	Weak References	✓	✗	✓	✓	✓
non-functional	Popularity	medium	medium	low	high	low
	Maturity	medium	medium	low	high	low
	Documentation	medium	medium	low	high	low

⁵<http://activemq.apache.org>

⁶<http://www.rabbitmq.com>

⁷<http://zeromq.org>

⁸<https://github.com/bennidi/mbassador>

⁹<https://code.google.com/p/guava-libraries/wiki/EventBusExplained>

¹⁰<https://code.google.com/p/simpleeventbus/>

¹¹<http://eventbus.org/>

¹²<https://github.com/mycila/pubsub>

Table 6.2: Feature comparison of Java PubSub libraries

7 Summary and Conclusion

Bibliography

- [1] *About AWS*. Amazon Web Services, Inc. URL: <http://aws.amazon.com/about-aws/>.
- [2] *Amazon EC2*. Amazon Web Services, Inc. URL: <http://aws.amazon.com/ec2/>.
- [3] *Amazon EC2 Instances*. Amazon Web Services, Inc. URL: <http://aws.amazon.com/ec2/instance-types/>.
- [4] *AWS Elastic Beanstalk (Beta)*. Amazon Web Services, Inc. URL: <http://aws.amazon.com/elasticbeanstalk/>.
- [5] Ralf Biedert. *Java Simple Plugin Framework F.A.Q.* URL: <https://code.google.com/p/jspf/wiki/FAQ>.
- [6] Tobias Binz et al. "OpenTOSCA -- A Runtime for TOSCA-based Cloud Applications". In: International Conference on Service-Oriented Computing. LNCS. Springer, 2013. URL: http://www.iaas.uni-stuttgart.de/RUS-data/INPROC-2013-45%20-%20OpenTOSCA_A_Runtime_for_TOSCA-based_Cloud_Applications.pdf.
- [7] Werner Buchholz. "The System Design of the IBM Type 701 Computer". In: Institute of Radio Engineers (IRE). Vol. 41. 10. 1953, pp. 1262–1275. URL: <http://ftp.cs.duke.edu/~xwy/publications/cloudcmp-imc10.pdf>.
- [8] *Definition of Provisioning*. ATIS Telecom Glossary. URL: <http://www.atis.org/glossary/definition.aspx?id=2474>.
- [9] Patrick Th. Eugster et al. "The many faces of publish/subscribe". In: *ACM Computing Surveys (CSUR)* 35.2 (June 2003), pp. 114–131. URL: core.kmi.open.ac.uk/download/pdf/12642066.pdf.
- [10] *Excellence Initiative at a Glance*. Tech. rep. German Research Foundation, Aug. 2013. URL: http://www.dfg.de/download/pdf/dfg_im_profil/geschaeftsstelle/publikationen/exin_broschuere_1307_en.pdf.
- [11] *Flexibility of Simulation Workflows*. SRC Simulation Technology. URL: http://www.iaas.uni-stuttgart.de/forschung/projects/simtech/project_flexibility.php.
- [12] *Getting Started with the AWS Management Console*. Amazon Web Services, Inc. URL: <http://docs.aws.amazon.com/awsconsolehelpdocs/latest/gsg/getting-started.html>.

Bibliography

- [13] James Joyce. *Ulysses*. Project Gutenberg. URL: <http://www.gutenberg.org/files/4300/4300-h/4300-h.htm>.
- [14] Johannes Kirschnick et al. "Toward an Architecture for the Automated Provisioning of Cloud Services". In: *Communications Magazine* 48.12 (2010), pp. 124–131. URL: <http://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=5673082>.
- [15] Johannes Kirschnick and et al. "Toward an architecture for the automated provisioning of cloud services". In: *Communications Magazine, IEEE* 48.12 (Dec. 2010), pp. 124–131. URL: <http://jmalcaraz.com/wp-content/uploads/papers/AlcarazCalero-2010-CommMag-Preprint.pdf>.
- [16] Ang Li et al. "CloudCmp: Comparing Public Cloud Providers". In: 10th ACM SIGCOMM conference on Internet measurement. 2010, pp. 1–14. URL: <http://ftp.cs.duke.edu/~xwy/publications/cloudcmp-imc10.pdf>.
- [17] Peter Mell and Timothy Grance. *The NIST Definition of Cloud Computing*. Tech. rep. National Institute of Standards and Technology, Sept. 2011. URL: <http://csrc.nist.gov/publications/nistpubs/800-145/SP800-145.pdf>.
- [18] *OSGi Service Platform core Specification*. Tech. rep. OSGi Alliance, Apr. 2011. URL: <http://www.osgi.org/download/r4v43/osgi.core-4.3.0.pdf>.
- [19] *PN 8: Integrated data management, workflow and visualisation to enable an integrative systems science*. SRC Simulation Technology. URL: <http://www.simtech.uni-stuttgart.de/forschung/pn/pn8/index.en.html>.
- [20] Javier Rojas. *The art of the bootstrap*. URL: <http://venturebeat.com/2008/11/20/the-art-of-the-bootstrap/>.
- [21] Valeri Schneider. "Dynamische Provisionierung von Web Services für Simulationsworkflows". Diploma. University Stuttgart, 2013. URL: ftp://ftp.informatik.uni-stuttgart.de/pub/library/medoc.ustuttgart_fi/DIP-3473/DIP-3473.pdf.
- [22] *ServiceLoader*. Oracle Corporation. URL: <http://docs.oracle.com/javase/6/docs/api/java/util/ServiceLoader.html>.
- [23] *SLA (Service-level Agreement)*. Gartner, Inc. URL: <http://www.gartner.com/it-glossary/sla-service-level-agreement/>.
- [24] *Tools for Amazon Web Services*. Amazon Web Services, Inc. URL: <https://aws.amazon.com/tools/>.
- [25] *Topology and Orchestration Specification for Cloud Applications*. Committee Specification 01. OASIS, 2013. URL: <http://docs.oasis-open.org/tosca/TOSCA/v1.0/cs01/TOSCA-v1.0-cs01.pdf>.

Bibliography

- [26] Luis M. Vaquero et al. "A Break in the Clouds: Towards a Cloud Definition". In: *ACM SIGCOMM Computer Communication Review* 39.1 (Jan. 2009), pp. 50–55. URL: <http://www.sigcomm.org/sites/default/files/ccr/papers/2009/January/1496091-1496100.pdf>.
- [27] Karolina Vukojevic-Haupt, Dimka Karastoyanova, and Frank Leymann. "On-demand Provisioning of Infrastructure, Middleware and Services for Simulation Workflows". In: *Proceedings of the 6th IEEE International Conference on Service Oriented Computing & Applications (SOCA 2013)*. 2013. URL: ???.
- [28] Benjamin Zimmer. *figurative 'bootstraps' (1834)*. The American Dialect Society Mailing List. URL: <http://listserv.linguistlist.org/cgi-bin/wa?A2=ind0508B&L=ADS-L&D=0&P=14972>.

All links were last visited on May 6, 2014.

Declaration of Authorship

I hereby certify that the diploma thesis entitled

Bootstrapping Provisioning Engines for On-demand Provisioning in Cloud Environments

is entirely my own work except where otherwise indicated. Passages and ideas from other sources have been clearly indicated. To date, neither this diploma thesis nor essential parts thereof were subject of an examination procedure. Until now, I don't have published this diploma thesis or parts thereof. The electronic copy is identical to the submitted copy.

Stuttgart, May 6, 2014,

.....
(Lukas Reinfurt)